

## PRACTICAL 1

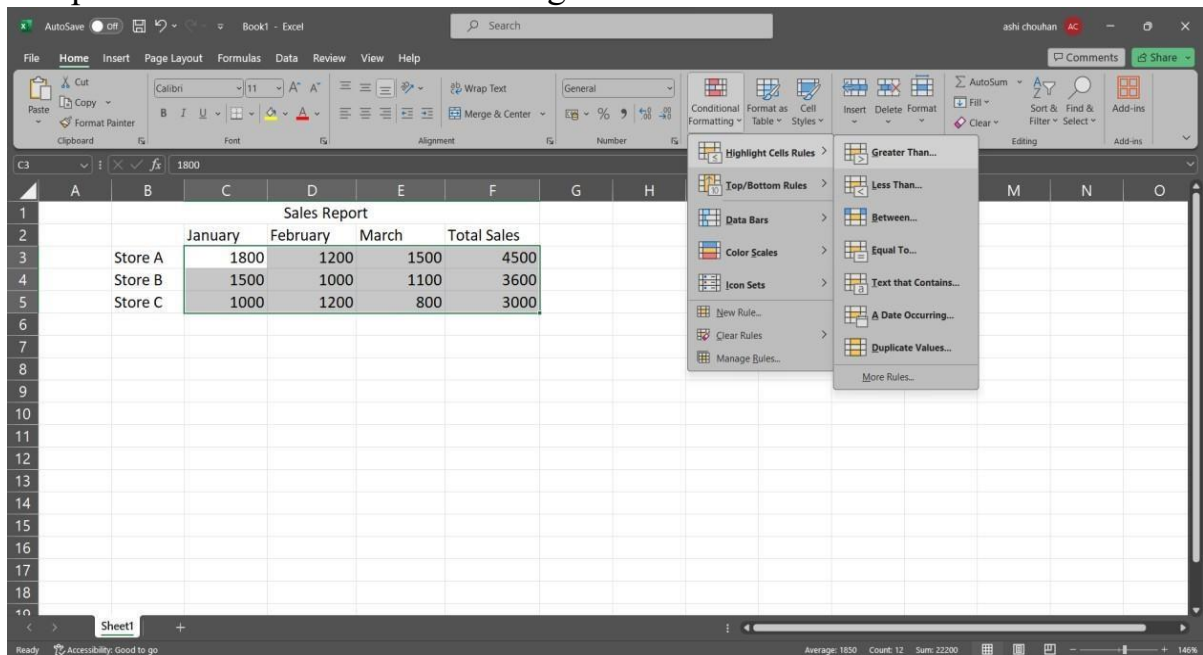
### Introduction to Excel

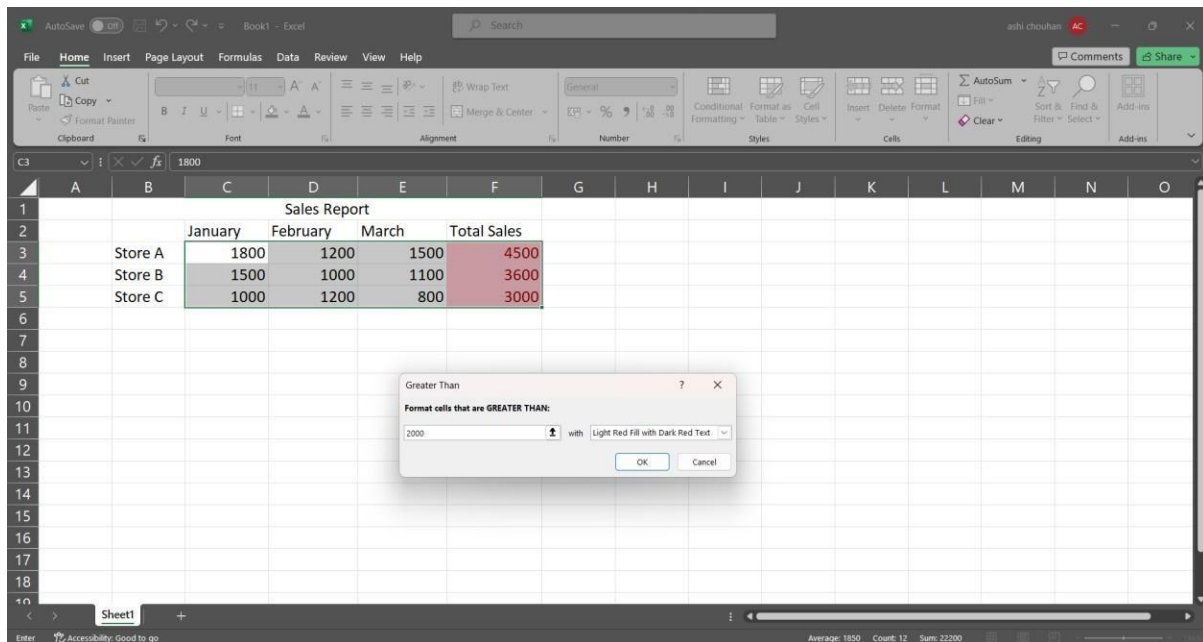
A. Perform conditional formatting on a dataset using various criteria.

	A	B	C	D	E	F	G	H
1				Sales Report				
2			January	February	March	Total Sales		
3		Store A	1800	1200	1500	4500		
4		Store B	1500	1000	1100	3600		
5		Store C	1000	1200	800	3000		
6								
7								
8								
9								
10								
11								
12								
13								

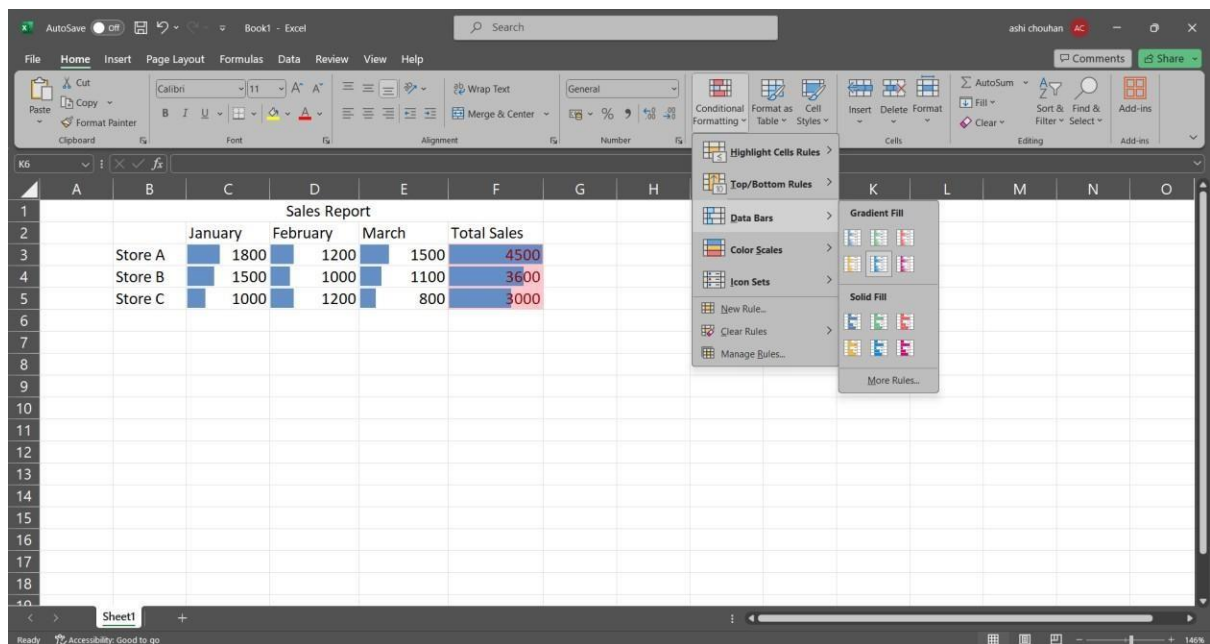
### Steps

Step 1: Go to conditional formatting > Greater Than





Step 3: Go to Data Bars > Solid Fill in conditional formatting.

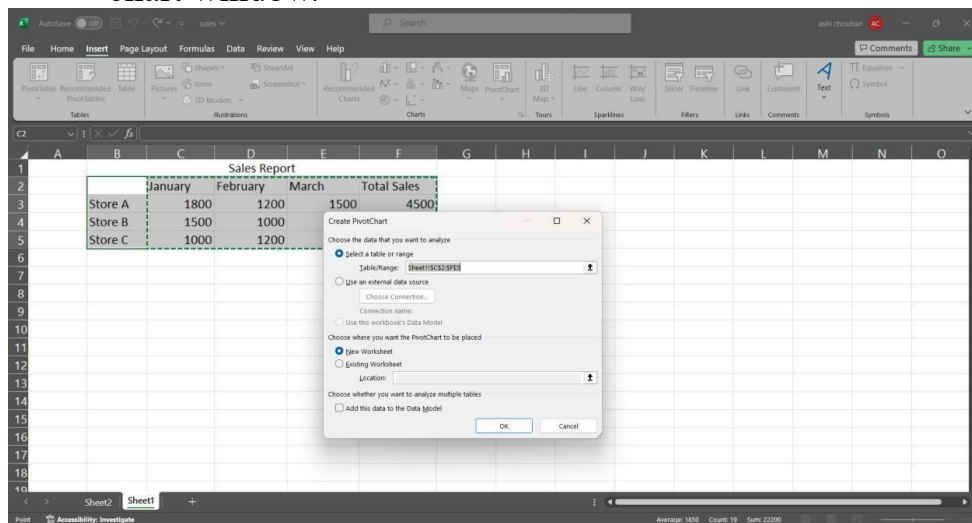


A. Create a pivot table to analyse and summarize data.

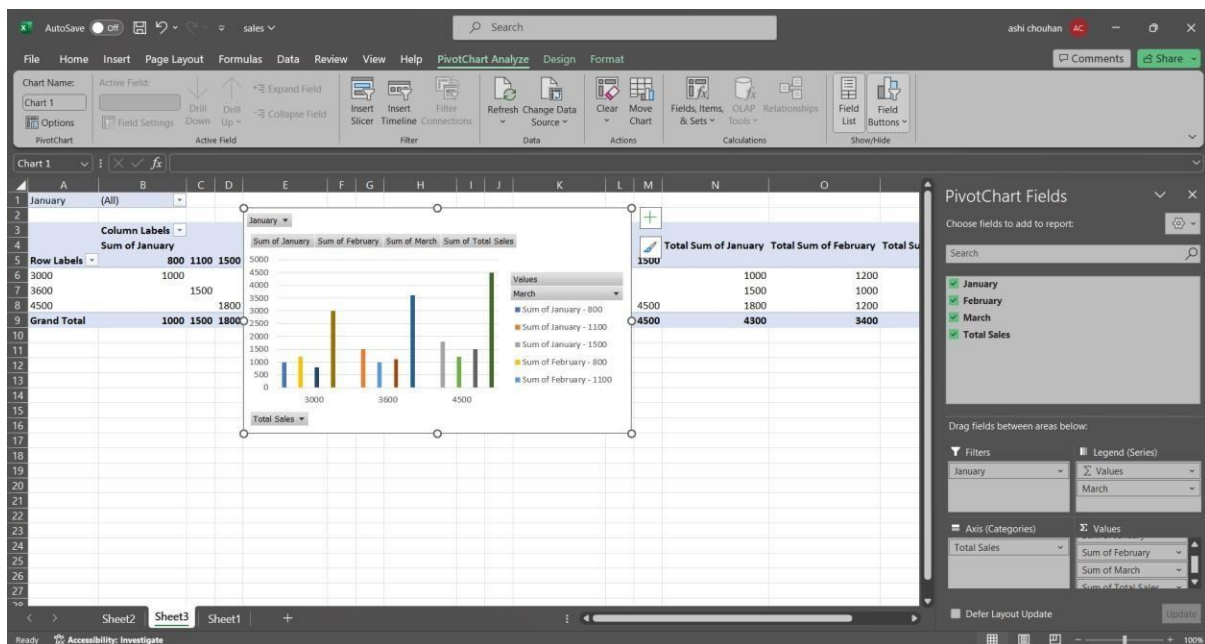
### Steps

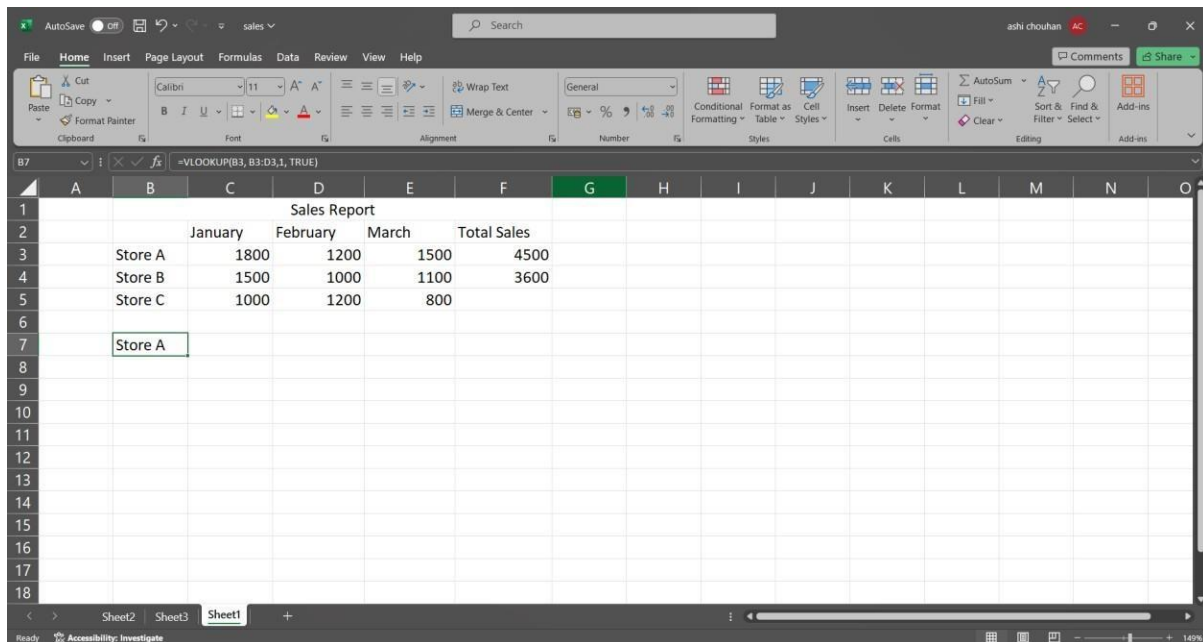
Step 1: select the entire table and go to Insert tab PivotChart > Pivotchart .

Step 2: Select “New worksheet” in the create pivot chart window.



Step 3: Select and drag attributes in the below boxes.

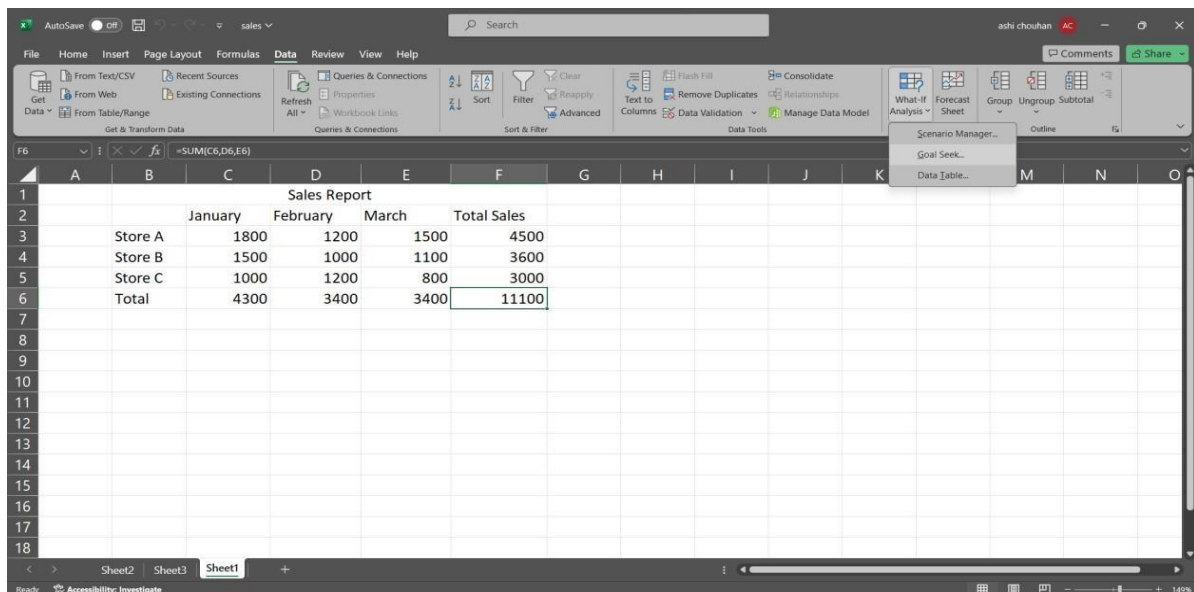




B. Perform what-if analysis using Goal Seek to determine input values for desired output.

Steps-

Step 1: In the Data tab go to the what if analysis>Goal seek



Step 2: Fill the information in the window accordingly and click ok

AutoSave | sales | Search | File | Home | Insert | Page Layout | Formulas | Data | Review | View | Help | Comments | Share

Get Data: From Text/CSV, Recent Sources, From Web, Existing Connections, From Table/Range, Workbook Links

Queries & Connections: Refresh All, Properties, Manage Data Model

Sort & Filter: Sort, Filter, Clear, Reapply, Advanced

Data Tools: Flash Fill, Remove Duplicates, Relationships, What-If Analysis, Data Validation, Manage Data Model

Forecast: Forecast Sheet

Outline: Group, Ungroup, Subtotal

F6: =SUM(C6,D6,E6)

		January	February	March	Total Sales
Store A		1800	1200	1500	4500
Store B		1500	1000	1100	3600
Store C		1000	1200	800	3000
Total		4300	3400	3400	11100

Goal Seek

Set cell: F6

To value: 12000

By changing cell: E6

OK Cancel

Sheet2 | Sheet3 | Sheet1

Enter | Accessibility: Investigate

AutoSave | sales | Search | File | Home | Insert | Page Layout | Formulas | Data | Review | View | Help | Comments | Share

Get Data: From Text/CSV, Recent Sources, From Web, Existing Connections, From Table/Range, Workbook Links

Queries & Connections: Refresh All, Properties, Manage Data Model

Sort & Filter: Sort, Filter, Clear, Reapply, Advanced

Data Tools: Flash Fill, Remove Duplicates, Relationships, What-If Analysis, Data Validation, Manage Data Model

Forecast: Forecast Sheet

Outline: Group, Ungroup, Subtotal

F6: =SUM(C6,D6,E6)

		January	February	March	Total Sales
Store A		1800	1200	1500	4500
Store B		1500	1000	1100	3600
Store C		1000	1200	109700	111900
Total		4300	3400	112300	120000

Goal Seek Status

Goal Seeking with Cell F6 found a solution.

Target value: 120000

Current value: 120000

Step | Pause | OK | Cancel

Sheet2 | Sheet3 | Sheet1

Ready | Accessibility: Investigate

AutoSave OFF sales

File Home Insert Page Layout Formulas Data Review View Help

Get Data From Text/CSV Recent Sources From Web Existing Connections From Table/Range

Refresh All Queries & Connections Properties Workbook Links

Sort Filter Clear Reapply Advanced

Text to Columns Remove Duplicates Data Validation Manage Data Model

Consolidate Relationships What-If Analysis Forecast Group Ungroup Subtotal

Comments Share

F6 =SUM(C6,D6,E6)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1				Sales Report											
2			January	February	March	Total Sales									
3		Store A	1800	1200	1500	4500									
4		Store B	1500	1000	1100	3600									
5		Store C	1000	1200	109700	111900									
6		Total	4300	3400	112300	120000									
7															
8															
9															
10															
11															
12															
13															
14															
15															
16															
17															
18															

Sheet2 Sheet3 Sheet1

Ready Accessibility: Investigate

## PRACTICAL 2

### Aim: Data Frames and Basic Data Pre-processing

#### **A. Read data from CSV and JSON files into a data frame.**

##### Code:

```
#Read data from CSV and JSON files into a data frame.
# Read data from a csv file
import pandas as pd
df = pd.read_csv('Student_Marks.csv')
print("Our dataset",df)
# Reading data from a JSON file
import pandas as pd
df = pd.read_json('dataset.json')
print("Our dataset",df)
```

##### Output:

```
= RESTART: C:\Users\Neeraj\Desktop\Tyics\sem 6\Data Science\practical\practical no 2 a.py
Our dataset  Unnamed: 0  Gender  DOB  Maths  ...  Biology  Economics  History  Civics
0      John      M  05-04-1988   55  ...      21      52      89      65
1    Suresh      M  04-05-1987   75  ...      90      61      58       2
2    Ramesh      M  25-05-1989   25  ...      95      87      56      74
3   Jessica      F  12-08-1990   78  ...      54      89      75      45
4   Jennifer      F  02-09-1989   58  ...      96      77      83      53
5     Annu      F  05-04-1988   45  ...      55      89      87      52
6    pooja      F  04-05-1987   55  ...      75      58      64      61
7    Ritesh      M  25-05-1989   54  ...      25      56      76      87
8    Farha      F  12-08-1990   55  ...      78      75      63      89
9    Mukesh      M  02-09-1989   96  ...      58      83      46      77

[10 rows x 11 columns]
Our dataset  fruit  size  color
0   Apple  Large   Red
1  Banana Medium Yellow
2   Orange Small  Orange
```

#### **B. Perform basic data pre-processing tasks such as handling missing values and outliers.**

##### Code:

```
#Perform basic data pre-processing tasks such as handling missing values and outliers.
# Replacing NA values using fillna()
import pandas as pd
df = pd.read_csv('titanic.csv')
print(df)
df.head(10)
print("Dataset after filling NA values with 0:")
df2 = df.fillna(value=0)
print(df2)
```

##### Output:

```
= RESTART: C:\Users\Neeraj\Desktop\Tyics\sem 6\Data Science\practical\practical no 2 b.py
```

	PassengerId	Survived	Pclass	...	Fare	Cabin	Embarked
0	1	0	3	...	7.2500	NaN	S
1	2	1	1	...	71.2833	C85	C
2	3	1	3	...	7.9250	NaN	S
3	4	1	1	...	53.1000	C123	S
4	5	0	3	...	8.0500	NaN	S
...	...	...	...	...	...	...	...
886	887	0	2	...	13.0000	NaN	S
887	888	1	1	...	30.0000	B42	S
888	889	0	3	...	23.4500	NaN	S
889	890	1	1	...	30.0000	C148	C
890	891	0	3	...	7.7500	NaN	Q

[891 rows x 12 columns]

Dataset after filling NA values with 0:

	PassengerId	Survived	Pclass	...	Fare	Cabin	Embarked
0	1	0	3	...	7.2500	0	S
1	2	1	1	...	71.2833	C85	C
2	3	1	3	...	7.9250	0	S
3	4	1	1	...	53.1000	C123	S
4	5	0	3	...	8.0500	0	S
...	...	...	...	...	...	...	...
886	887	0	2	...	13.0000	0	S
887	888	1	1	...	30.0000	B42	S
888	889	0	3	...	23.4500	0	S
889	890	1	1	...	30.0000	C148	C
890	891	0	3	...	7.7500	0	Q

[891 rows x 12 columns]

### C. Manipulate and transform data using functions like filtering, sorting, and grouping.

#### Code:

```
import pandas as pd
# Load dataset
iris = pd.read_csv('Iris.csv')
# Check column names
print("Columns:", iris.columns)
# Filter 'setosa' species (correct column name is 'target')
setosa = iris[iris['target'] == 'Iris-setosa']
print(setosa.head())
# Sorting dataset
sorted_iris = iris.sort_values(by='sepal length (cm)', ascending=False)
print(sorted_iris.head())
# Group by 'target'
grouped_species = iris.groupby('target').mean(numeric_only=True)
print(grouped_species)
```

#### Output:



```
= RESTART: C:\Users\Neeraj\Desktop\Tyics\sem 6\Data Science\practical\practical no 2 c.py
Columns: Index(['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)',
               'petal width (cm)', 'target'],
              dtype='object')
   sepal length (cm)  sepal width (cm)  ...  petal width (cm)  target
0                5.1                3.5  ...                0.2  Iris-setosa
1                4.9                3.0  ...                0.2  Iris-setosa
2                4.7                3.2  ...                0.2  Iris-setosa
3                4.6                3.1  ...                0.2  Iris-setosa
4                5.0                3.6  ...                0.2  Iris-setosa

[5 rows x 5 columns]
   sepal length (cm)  sepal width (cm)  ...  petal width (cm)  target
131                7.9                3.8  ...                2.0  Iris-virginica
122                7.7                2.8  ...                2.0  Iris-virginica
118                7.7                2.6  ...                2.3  Iris-virginica
117                7.7                3.8  ...                2.2  Iris-virginica
135                7.7                3.0  ...                2.3  Iris-virginica

[5 rows x 5 columns]
   sepal length (cm)  ...  petal width (cm)
target              ...
Iris-setosa          5.006  ...          0.244
Iris-versicolor      5.936  ...          1.326
Iris-virginica       6.588  ...          2.026

[3 rows x 4 columns]
```

---

## PRACTICAL 3

### Aim: Feature Scaling and Dummification

**A. Apply feature-scaling techniques like standardization and normalization to numerical features.**

#### Code:

```
# Apply feature-scaling techniques like standardization and normalization to
numerical features.
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler, StandardScaler
# Load the dataset
df = pd.read_csv('wine.csv', header=None, usecols=[0, 1, 2], skiprows=1)
df.columns = ['classlabel', 'Alcohol', 'Malic Acid']
# Display the original dataframe
print("Original DataFrame:")
print(df.head()) # Print first few rows for better readability
# Apply Min-Max Scaling
minmax_scaler = MinMaxScaler()
df[['Alcohol', 'Malic Acid']] = minmax_scaler.fit_transform(df[['Alcohol', 'Malic Acid']])
print("\nDataFrame after Min-Max Scaling:")
print(df.head())
# Apply Standard Scaling (Z-score normalization)
standard_scaler = StandardScaler()
df[['Alcohol', 'Malic Acid']] = standard_scaler.fit_transform(df[['Alcohol', 'Malic Acid']])
print("\nDataFrame after Standard Scaling:")
print(df.head())
```

#### Output:

```
> = RESTART: C:\Users\Neeraj\Desktop\Tyics\sem 6\Data Science\practical\practical no 3 a.py
Original DataFrame:
   classlabel  Alcohol  Malic Acid
0           1    14.23      1.71
1           1    13.20      1.78
2           1    13.16      2.36
3           1    14.37      1.95
4           1    13.24      2.59

DataFrame after Min-Max Scaling:
   classlabel  Alcohol  Malic Acid
0           1  0.884298  0.000000
1           1  0.033058  0.079545
2           1  0.000000  0.738636
3           1  1.000000  0.272727
4           1  0.066116  1.000000

DataFrame after Standard Scaling:
   classlabel  Alcohol  Malic Acid
0           1  1.089979 -1.078368
1           1 -0.812866 -0.873243
2           1 -0.886763  0.826358
3           1  1.348618 -0.375084
4           1 -0.738969  1.500338
```

**B. Perform feature dummification to convert categorical variables into numerical representations.**

**Code:**

```
#Perform feature dummification to convert categorical variables into numerical
representations.
import pandas as pd
from sklearn.preprocessing import LabelEncoder
# Load dataset
iris = pd.read_csv("Iris.csv")
# Check column names
print("Columns in dataset:", iris.columns)
# Encode the 'target' column instead of 'Species'
le = LabelEncoder()
iris['code'] = le.fit_transform(iris['target'])
# Print updated DataFrame
print(iris.head())
```

**Output:**

```
> C:\Users\Neeraj\OneDrive\Documents> python Iris.py
= RESTART: C:\Users\Neeraj\Desktop\Tyca\sem 6\Data Science\practical\practical no 3 b.py
Columns in dataset: Index(['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)',
      'petal width (cm)', 'target'],
      dtype='object')
   sepal length (cm)  sepal width (cm)  ...      target  code
0              5.1              3.5  ...  Iris-setosa    0
1              4.9              3.0  ...  Iris-setosa    0
2              4.7              3.2  ...  Iris-setosa    0
3              4.6              3.1  ...  Iris-setosa    0
4              5.0              3.6  ...  Iris-setosa    0

[5 rows x 6 columns]
```

## PRACTICAL 4

### Aim: Hypothesis Testing

- A. Formulate null and alternative hypotheses for a given problem.**
- B. Conduct a hypothesis test using appropriate statistical tests (e.g., t-test, chisquare test).**
- C. Interpret the results and draw conclusions based on the test outcomes.**

### Code:

```
import numpy as np
from scipy import stats
import matplotlib.pyplot as plt
# Generate two samples for demonstration purposes
np.random.seed(42)
sample1 = np.random.normal(loc=10, scale=2, size=30)
sample2 = np.random.normal(loc=12, scale=2, size=30)
# Perform a two-sample t-test
t_statistic, p_value = stats.ttest_ind(sample1, sample2)
# Set the significance level
alpha = 0.05
# Print the results of the two-sample t-test
print("Results of Two-Sample t-test:")
print(f'T-statistic: {t_statistic}')
print(f'P-value: {p_value}')
print(f'Degrees of Freedom: {len(sample1) + len(sample2) - 2}')
# Plot the distributions of the two samples
plt.figure(figsize=(10, 6))
plt.hist(sample1, alpha=0.5, label='Sample 1', color='blue', bins=10)
plt.hist(sample2, alpha=0.5, label='Sample 2', color='orange', bins=10)
plt.axvline(np.mean(sample1), color='blue', linestyle='dashed', linewidth=2)
plt.axvline(np.mean(sample2), color='orange', linestyle='dashed', linewidth=2)
plt.title('Distributions of Sample 1 and Sample 2')
plt.xlabel('Values')
plt.ylabel('Frequency')
plt.legend()
# Highlight the critical region if the null hypothesis is rejected
if p_value < alpha:
    critical_region = np.linspace(min(sample1.min(), sample2.min()), max(sample1.max(),
sample2.max()), 1000)
    plt.fill_between(critical_region, 0, 5, color='red', alpha=0.3, label='Critical Region')
    plt.text(11, 5, f'T-statistic: {t_statistic:.2f}', ha='center', va='center', color='black',
backgroundcolor='white')
# Show the plot
plt.show()
# Draw Conclusions
if p_value < alpha:
```

```

if np.mean(sample1) > np.mean(sample2):
    print("Conclusion: There is significant evidence to reject the null hypothesis.")
    print("Interpretation: The mean of Sample 1 is significantly higher than that of Sample 2.")
else:
    print("Conclusion: There is significant evidence to reject the null hypothesis.")
    print("Interpretation: The mean of Sample 2 is significantly higher than that of Sample 1.")
else:
    print("Conclusion: Fail to reject the null hypothesis.")
    print("Interpretation: There is not enough evidence to claim a significant difference between the
means.")


```

### **Output:**

```

= RESTART: C:\Users\Neeraj\Desktop\Tyics\sem 6\Data Science\practical\practical no 4.py
Results of Two-Sample t-test:
T-statistic: -4.512913234547555
P-value: 3.176506547470154e-05
Degrees of Freedom: 58

```



```

#chi-test

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb
import warnings
from scipy import stats
# Ignore warnings
warnings.filterwarnings('ignore')
# Load the dataset
df = sb.load_dataset('mpg')
# Display the dataframe and describe columns
print(df)
print(df['horsepower'].describe())
print(df['model_year'].describe())
# Define bins and categorize 'horsepower'
bins = [0, 75, 150, 240]
df['horsepower_new'] = pd.cut(df['horsepower'], bins=bins, labels=['l', 'm', 'h'])
# Display the new 'horsepower' categories
c = df['horsepower_new']
print(c)
# Define bins for 'model_year' and categorize
ybins = [69, 72, 74, 84]
labels = ['t1', 't2', 't3']
df['modelyear_new'] = pd.cut(df['model_year'], bins=ybins, labels=labels)
# Display the new 'model_year' categories
newyear = df['modelyear_new']
print(newyear)
# Perform chi-squared test of independence
df_chi = pd.crosstab(df['horsepower_new'], df['modelyear_new'])
print(df_chi)
# Perform chi-squared test
chi2_stat, p_value, dof, expected = stats.chi2_contingency(df_chi)
print(f"Chi2 Statistic: {chi2_stat}")
print(f"P-value: {p_value}")
print(f"Degrees of Freedom: {dof}")
print(f"Expected Frequencies: \n{expected}")

```

```
= RESTART: C:\Users\Neeraj\Desktop\Tyics\sem 6\Data Science\horsepower.py
```

	mpg	cylinders	...	origin	name
0	18.0	8	...	usa	chevrolet chevelle malibu
1	15.0	8	...	usa	buick skylark 320
2	18.0	8	...	usa	plymouth satellite
3	16.0	8	...	usa	amc rebel sst
4	17.0	8	...	usa	ford torino
..	...	...	...	...	...
393	27.0	4	...	usa	ford mustang gl
394	44.0	4	...	europa	vw pickup
395	32.0	4	...	usa	dodge rampage
396	28.0	4	...	usa	ford ranger
397	31.0	4	...	usa	chevy s-10

```
[398 rows x 9 columns]
```

```
count    392.000000
mean     104.469388
std       38.491160
min       46.000000
25%       75.000000
50%       93.500000
75%      126.000000
max      230.000000
```

```
Name: horsepower, dtype: float64
```

```
count    398.000000
mean       76.010050
std         3.697627
min        70.000000
25%        73.000000
50%        76.000000
75%        79.000000
max        82.000000
```

```
Name: model_year, dtype: float64
```

```
0      m
1      h
2      m
3      m
4      m
```

```
..
393    m
394    l
395    m
396    m
397    m
```

```
Name: horsepower_new, Length: 398, dtype: category
```

```
Categories (3, object): ['l' < 'm' < 'h']
```

```
0      t1
1      t1
2      t1
3      t1
4      t1
..
```

```

393     t3
394     t3
395     t3
396     t3
397     t3
Name: modelyear_new, Length: 398, dtype: category
Categories (3, object): ['t1' < 't2' < 't3']
modelyear_new  t1  t2  t3
horsepower_new
l               9  14  76
m              49  41 158
h              26  11   8
Chi2 Statistic: 54.95485392447537
P-value: 3.320518009555984e-11
Degrees of Freedom: 4
Expected Frequencies:
[[ 21.21428571  16.66836735  61.11734694]
 [ 53.14285714  41.75510204 153.10204082]
 [   9.64285714   7.57653061  27.78061224]]

```



## PRACTICAL 5

### Aim: ANOVA (Analysis of Variance)

- Perform one-way ANOVA to compare means across multiple groups.
- Conduct post-hoc tests to identify significant differences between group means.

### Code:

```
import pandas as pd
import scipy.stats as stats
from statsmodels.stats.multicomp import pairwise_tukeyhsd
# Data for each group
group1 = [23, 25, 29, 34, 30]
group2 = [19, 20, 22, 24, 25]
group3 = [15, 18, 20, 21, 17]
group4 = [28, 24, 26, 30, 29]
# Combine data into a DataFrame
data = pd.DataFrame({
    'value': group1 + group2 + group3 + group4,
    'group': ['Group1'] * len(group1) + ['Group2'] * len(group2) + ['Group3'] * len(group3) + ['Group4']
    * len(group4)
})
# Perform one-way ANOVA
f_statistics, p_value = stats.f_oneway(group1, group2, group3, group4)
print("One-way ANOVA:")
print("F-statistic:", f_statistics)
print("p-value:", p_value)
# Perform Tukey-Kramer post-hoc test if ANOVA is significant
if p_value < 0.05:
    tukey_results = pairwise_tukeyhsd(data['value'], data['group'])
    print("\nTukey-Kramer post-hoc test:")
    print(tukey_results)
else:
    print("\nNo significant difference found in ANOVA, so Tukey-Kramer test is not performed.")
```

### Output:

```
= RESTART: C:\Users\Neeraj\Desktop\Tyces\sem 6\Data Science\practical\practical no 5.py
One-way ANOVA:
F-statistic: 12.139872842870115
p-value: 0.00021465200901629603

Tukey-Kramer post-hoc test:
Multiple Comparison of Means - Tukey HSD, FWER=0.05
=====
group1 group2 meandiff p-adj lower upper reject
-----
Group1 Group2 -6.2 0.024 -11.6809 -0.7191 True
Group1 Group3 -10.0 0.0004 -15.4809 -4.5191 True
Group1 Group4 -0.8 0.9747 -6.2809 4.6809 False
Group2 Group3 -3.8 0.2348 -9.2809 1.6809 False
Group2 Group4 5.4 0.0542 -0.0809 10.8809 False
Group3 Group4 9.2 0.001 3.7191 14.6809 True
=====
```

## **PRACTICAL 6**

### **Aim: Regression and Its Types**

- **Implement simple linear regression using a dataset.**
- **Explore and interpret the regression model coefficients and goodness-of-fit measures.**

#### **Code:**

```
import numpy as np
import pandas as pd
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
# Load California housing dataset
print("Loading the California housing dataset...")
housing = fetch_california_housing()
housing_df = pd.DataFrame(housing.data, columns=housing.feature_names)
housing_df['PRICE'] = housing.target
# Check first few rows of the dataframe to make sure data is loaded
print(housing_df.head())
# Define features (X) and target (y)
X = housing_df[['AveRooms']]
y = housing_df['PRICE']
# Split the data into training and testing sets
print("Splitting the data into training and testing sets...")
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Train the model
print("Training the model...")
model = LinearRegression()
model.fit(X_train, y_train)
print("Model training complete.")
# Make predictions and evaluate the model
print("Making predictions and evaluating the model...")
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
# Output the results
print("Mean Squared Error:", mse)
print("R-squared:", r2)
print("Intercept:", model.intercept_)
print("Coefficient:", model.coef_)
```

#### **Output:**

```
= RESTART: C:\Users\Neeraj\Desktop\Tyics\sem 6\Data Science\practical\practical no 6.py
Loading the California housing dataset...
MedInc HouseAge AveRooms AveBedrms ... AveOccup Latitude Longitude PRICE
0 8.3252 41.0 6.984127 1.023810 ... 2.555556 37.88 -122.23 4.526
1 8.3014 21.0 6.238137 0.971880 ... 2.109842 37.86 -122.22 3.585
2 7.2574 52.0 8.288136 1.073446 ... 2.802260 37.85 -122.24 3.521
3 5.6431 52.0 5.817352 1.073059 ... 2.547945 37.85 -122.25 3.413
4 3.8462 52.0 6.281853 1.081081 ... 2.181467 37.85 -122.25 3.422

[5 rows x 9 columns]
Splitting the data into training and testing sets...
Training the model...
Model training complete.
Making predictions and evaluating the model...
Mean Squared Error: 1.2923314440807299
R-squared: 0.013795337532284901
Intercept: 1.654762268596842
Coefficient: [0.07675559]
```

- **Extend the analysis to multiple linear regression and assess the impact of additional predictors.**

#### **Code:**

```
import numpy as np
import pandas as pd
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
# Load California housing dataset
housing = fetch_california_housing()
housing_df = pd.DataFrame(housing.data, columns=housing.feature_names)
housing_df['PRICE'] = housing.target
# Define features (X) and target (y)
X = housing_df.drop('PRICE', axis=1) # Drop 'PRICE' column from the features
y = housing_df['PRICE'] # 'PRICE' is the target variable
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)
# Make predictions
y_pred = model.predict(X_test)
# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
# Output the results
print("Mean Squared Error:", mse)
print("R-squared:", r2)
```

```
print("Intercept:", model.intercept_)  
print("Coefficients:", model.coef_)
```

**Output:**

```
= RESTART: C:/Users/Neeraj/Desktop/Tyacs/sem 6/Data Science/practical/6 b.py  
Mean Squared Error: 0.555891598695244  
R-squared: 0.5757877060324511  
Intercept: -37.023277706064064  
Coefficients: [ 4.48674910e-01  9.72425752e-03 -1.23323343e-01  7.83144907e-01  
-2.02962058e-06 -3.52631849e-03 -4.19792487e-01 -4.33708065e-01]  
|
```

## PRACTICAL 7

### Aim: Logistic Regression and Decision Tree

- Build a logistic regression model to predict a binary outcome.
- Evaluate the model's performance using classification metrics (e.g., accuracy, precision, recall).
- Construct a decision tree model and interpret the decision rules for classification.

### Code:

```
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, classification_report
# Load the Iris dataset and create a binary classification problem
iris = load_iris()
iris_df = pd.DataFrame(data=np.c_[iris['data'], iris['target']],
                      columns=iris['feature_names'] + ['target'])
# Filter the dataset to create a binary classification problem (target != 2)
binary_df = iris_df[iris_df['target'] != 2]
# Define features (X) and target (y)
X = binary_df.drop('target', axis=1)
y = binary_df['target']
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Train a logistic regression model and evaluate its performance
logistic_model = LogisticRegression()
logistic_model.fit(X_train, y_train)
y_pred_logistic = logistic_model.predict(X_test)
print("Logistic Regression Metrics")
print("Accuracy: ", accuracy_score(y_test, y_pred_logistic))
print("Precision:", precision_score(y_test, y_pred_logistic))
print("Recall: ", recall_score(y_test, y_pred_logistic))
print("\nClassification Report:")
print(classification_report(y_test, y_pred_logistic))
# Train a decision tree model and evaluate its performance
decision_tree_model = DecisionTreeClassifier()
decision_tree_model.fit(X_train, y_train)
y_pred_tree = decision_tree_model.predict(X_test)
print("\nDecision Tree Metrics")
print("Accuracy: ", accuracy_score(y_test, y_pred_tree))
print("Precision:", precision_score(y_test, y_pred_tree))
print("Recall: ", recall_score(y_test, y_pred_tree))
```

```
print("\nClassification Report:")
print(classification_report(y_test, y_pred_tree))
```

**Output:**

```
= RESTART: C:/Users/Neeraj/Desktop/Tycs/sem 6/Data Science/practical/practical no 7.py
Logistic Regression Metrics
Accuracy: 1.0
Precision: 1.0
Recall: 1.0

Classification Report:
              precision    recall  f1-score   support

    0.0         1.00      1.00      1.00        12
    1.0         1.00      1.00      1.00         8

   accuracy          1.00      1.00      1.00        20
  macro avg          1.00      1.00      1.00        20
weighted avg          1.00      1.00      1.00        20

Decision Tree Metrics
Accuracy: 1.0
Precision: 1.0
Recall: 1.0

Classification Report:
              precision    recall  f1-score   support

    0.0         1.00      1.00      1.00        12
    1.0         1.00      1.00      1.00         8

   accuracy          1.00      1.00      1.00        20
  macro avg          1.00      1.00      1.00        20
weighted avg          1.00      1.00      1.00        20
```

## **PRACTICAL 8**

### **Aim: K-Means Clustering**

- Apply the K-Means algorithm to group similar data points into clusters.
- Determine the optimal number of clusters using elbow method or silhouette analysis.
- Visualize the clustering results and analyze the cluster characteristics.

### **Code:**

```
import pandas as pd

from sklearn.preprocessing import MinMaxScaler
from sklearn.cluster import KMeans

import matplotlib.pyplot as plt

# Load the dataset
data = pd.read_csv("Wholesale.csv")

# Display the first few rows of the data
print(data.head())

# Define categorical and continuous features
categorical_features = ['Channel', 'Region']
continuous_features = ['Fresh', 'Milk', 'Grocery', 'Frozen', 'Detergents_Paper', 'Delicassen']

# Display summary statistics of the continuous features
print(data[continuous_features].describe())

# One-hot encoding for categorical features
for col in categorical_features:
    dummies = pd.get_dummies(data[col], prefix=col)
    data = pd.concat([data, dummies], axis=1)
    data.drop(col, axis=1, inplace=True)

# Display the data after encoding
print(data.head())

# Apply Min-Max Scaling to the data
mms = MinMaxScaler()
data_transformed = mms.fit_transform(data)


# Elbow method to determine the optimal number of clusters
```

```

sum_of_squared_distances = []
K = range(1, 15)
for k in K:
    km = KMeans(n_clusters=k)
    km.fit(data_transformed)
    sum_of_squared_distances.append(km.inertia_)
# Plotting the elbow graph
plt.plot(K, sum_of_squared_distances, 'bx-')
plt.xlabel('k (Number of Clusters)')
plt.ylabel('Sum of Squared Distances (Inertia)')
plt.title('Elbow Method for Optimal k')
plt.show()

```

**Output:**





## **PRACTICAL 9**

### **Aim: Principal Component Analysis (PCA)**


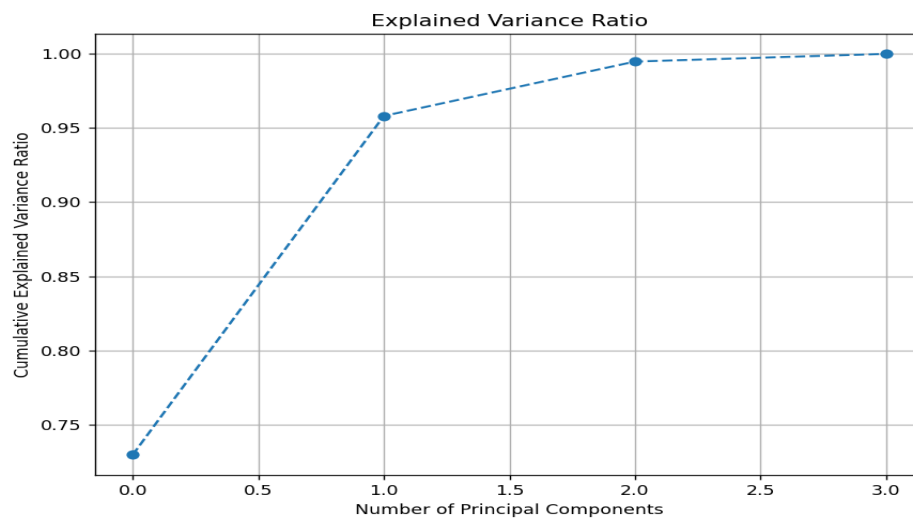
- Perform PCA on a dataset to reduce dimensionality.
- Evaluate the explained variance and select the appropriate number of principal components.
- Visualize the data in the reduced-dimensional space.

### **Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
# Load the Iris dataset
iris = load_iris()
iris_df = pd.DataFrame(data=np.c_[iris['data'], iris['target']],
                       columns=iris['feature_names'] + ['target'])
# Features and target variables
X = iris_df.drop('target', axis=1)
y = iris_df['target']
# Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
# Apply PCA
pca = PCA()
X_pca = pca.fit_transform(X_scaled)
# Explained variance ratio
explained_variance_ratio = pca.explained_variance_ratio_
# Plot the cumulative explained variance ratio
plt.figure(figsize=(8, 6))
plt.plot(np.cumsum(explained_variance_ratio), marker='o', linestyle='--')
plt.title('Explained Variance Ratio')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance Ratio')
plt.grid(True)
plt.show()
# Find the number of components that explain 95% of the variance
cumulative_variance_ratio = np.cumsum(explained_variance_ratio)
n_components = np.argmax(cumulative_variance_ratio >= 0.95) + 1
print(f'Number of principal components to explain 95% variance: {n_components}')
# Reduce the data to the chosen number of components
pca = PCA(n_components=n_components)
X_reduced = pca.fit_transform(X_scaled)
# Plot the reduced data in 2D space
```

```
plt.figure(figsize=(8, 6))
plt.scatter(X_reduced[:, 0], X_reduced[:, 1], c=y, cmap='viridis', s=50, alpha=0.5)
plt.title('Data in Reduced-dimensional Space')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Target')
plt.show()
```

**Output:**



```
= RESTART: C:/Users/Neeraj/Desktop/Tycs/sem 6/Data Science/practical/practical no 9.py
Number of principal components to explain 95% variance: 2
```

## **PRACTICAL 10**

### **Aim: Data Visualization and Storytelling**



- Create meaningful visualizations using data visualization tools
- Combine multiple visualizations to tell a compelling data story.
- Present the findings and insights in a clear and concise manner.


### **Code:**

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
# Generate random data
np.random.seed(42) # Set a seed for reproducibility
# Create a DataFrame with random data
data = pd.DataFrame({
    'variable1': np.random.normal(0, 1, 1000),
    'variable2': np.random.normal(2, 2, 1000) + 0.5 * np.random.normal(0, 1, 1000),
    'variable3': np.random.normal(-1, 1.5, 1000),
    'category': pd.Series(np.random.choice(['A', 'B', 'C', 'D'], size=1000, p=[0.4, 0.3, 0.2, 0.1]),
dtype='category')
})
# Create a scatter plot to visualize the relationship between two variables
plt.figure(figsize=(10, 6))
plt.scatter(data['variable1'], data['variable2'], alpha=0.5)
plt.title('Relationship between Variable 1 and Variable 2', fontsize=16)
plt.xlabel('Variable 1', fontsize=14)
plt.ylabel('Variable 2', fontsize=14)
plt.show()
# Create a bar chart to visualize the distribution of a categorical variable
plt.figure(figsize=(10, 6))
sns.countplot(x='category', data=data)
plt.title('Distribution of Categories', fontsize=16)
plt.xlabel('Category', fontsize=14)
plt.ylabel('Count', fontsize=14)
plt.xticks(rotation=45)
plt.show()
# Create a heatmap to visualize the correlation between numerical variables
plt.figure(figsize=(10, 8))
numerical_cols = ['variable1', 'variable2', 'variable3']
```

```
sns.heatmap(data[numerical_cols].corr(), annot=True, cmap='coolwarm', fmt='.2f', linewidths=1)
plt.title('Correlation Heatmap', fontsize=16)
plt.show()
# Data Storytelling
print("Title: Exploring the Relationship between Variable 1 and Variable 2\n")
print("The scatter plot (Figure 1) shows the relationship between Variable 1 and Variable 2.")
print("\nFigure 1: Scatter Plot of Variable 1 and Variable 2")
print("\nTo better understand the distribution of the categorical variable 'category', we created a bar chart.")
print("\nFigure 2: Distribution of Categories")
print("\nAdditionally, we explored the correlation between numerical variables using a heatmap.")
print("\nFigure 3: Correlation Heatmap")
print("\nIn summary, the visualizations and analysis provide insights into the relationships between the variables, such as the correlation between the numerical variables and the distribution of categories.")
```

### **Output:**





```
= RESTART: C:/Users/Neeraj/Desktop/Tyacs/sem 6/Data Science/practical/practical no 10.py
Title: Exploring the Relationship between Variable 1 and Variable 2
```

```
The scatter plot (Figure 1) shows the relationship between Variable 1 and Variable 2.
```

```
Figure 1: Scatter Plot of Variable 1 and Variable 2
```

```
To better understand the distribution of the categorical variable 'category', we created a bar chart.
```

```
Figure 2: Distribution of Categories
```

```
Additionally, we explored the correlation between numerical variables using a heatmap.
```

```
Figure 3: Correlation Heatmap
```

```
In summary, the visualizations and analysis provide insights into the relationships between the variables, such as the correlation between the numerical variables and the distribution of categories.
```